

Let's first import the packages

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

Import and display dataset's head

```
In [2]: df = pd.read_csv('Exam_Score_Prediction.csv')
print(df.head())

  student_id  age  gender  course  study_hours  class_attendance \
0      200.99   17    male  diploma        2.78             92.9
1      200.99   23  other    bca        3.37             64.8
2      200.99   22    male    b.sc        7.88             76.8
3      200.99   20  other  diploma        0.67             48.4
4      200.99   20  female  diploma        0.89             71.6

  internet_access  sleep_hours  sleep_quality  study_method  facility_ratin
g \
0           yes          7.4         poor       coaching        low
w
1           yes          4.6        average     online videos     medium
m
2           yes          8.5         poor       coaching        high
h
3           yes          5.8        average     online videos        low
w
4           yes          9.8         poor       coaching        low
w

  exam_difficulty  exam_score
0            hard       58.9
1      moderate       54.8
2      moderate       90.3
3      moderate       29.7
4      moderate       43.7
```

Display Dataset information

```
In [3]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   student_id      20000 non-null   float64
 1   age              20000 non-null   int64  
 2   gender            20000 non-null   object  
 3   course            20000 non-null   object  
 4   study_hours       20000 non-null   float64
 5   class_attendance 20000 non-null   float64
 6   internet_access  20000 non-null   object  
 7   sleep_hours       20000 non-null   float64
 8   sleep_quality     20000 non-null   object  
 9   study_method      20000 non-null   object  
 10  facility_rating  20000 non-null   object  
 11  exam_difficulty  20000 non-null   object  
 12  exam_score        20000 non-null   float64
dtypes: float64(5), int64(1), object(7)
memory usage: 2.0+ MB
None
```

Check if there is any missing values

```
In [4]: print("\nMissing Values:\n", df.isna().sum())
df.describe(include='all')
```

```
Missing Values:
student_id      0
age              0
gender            0
course            0
study_hours       0
class_attendance 0
internet_access  0
sleep_hours       0
sleep_quality     0
study_method      0
facility_rating  0
exam_difficulty  0
exam_score        0
dtype: int64
```

Out[4]:

	student_id	age	gender	course	study_hours	class_attendance
count	20000.000000	20000.000000	20000	20000	20000.000000	20000.000000
unique	Nan	Nan	3	7	Nan	Nan
top	Nan	Nan	other	bca	Nan	Nan
freq	Nan	Nan	6726	2902	Nan	Nan
mean	10000.500000	20.473300	Nan	Nan	4.007604	70.017365
std	5770.211372	2.284458	Nan	Nan	2.308313	17.282262
min	200.990000	17.000000	Nan	Nan	0.080000	40.600000
25%	5000.750000	18.000000	Nan	Nan	2.000000	55.100000
50%	10000.500000	20.000000	Nan	Nan	4.040000	69.900000
75%	15000.250000	22.000000	Nan	Nan	6.000000	85.000000
max	19800.010000	24.000000	Nan	Nan	7.910000	99.400000



There is no any missing values, so we move further

Divide into categorial columns

```
In [5]: numeric_cols = ["age", "study_hours", "class_attendance", "sleep_hours",
categorical_cols = [
    "gender", "course", "internet_access",
    "sleep_quality", "study_method",
    "facility_rating", "exam_difficulty"
]
```

If there is any missing values then fill with mean value

```
In [6]: # Create imputers
num_imputer = SimpleImputer(strategy='mean')
cat_imputer = SimpleImputer(strategy='most_frequent')

df[numeric_cols] = num_imputer.fit_transform(df[numeric_cols])
df[categorical_cols] = cat_imputer.fit_transform(df[categorical_cols])

df.head()
```

Out[6]:

	student_id	age	gender	course	study_hours	class_attendance	internet_access
0	200.99	17.0	male	diploma	2.78	92.9	yes
1	200.99	23.0	other	bca	3.37	64.8	yes
2	200.99	22.0	male	b.sc	7.88	76.8	yes
3	200.99	20.0	other	diploma	0.67	48.4	yes
4	200.99	20.0	female	diploma	0.89	71.6	yes

Encode categorical features (One-Hot Encoding)

In [7]:

```
df_encoded = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
df_encoded.head()
```

Out[7]:

	student_id	age	study_hours	class_attendance	sleep_hours	exam_score	gender_r
0	200.99	17.0	2.78	92.9	7.4	58.9	.
1	200.99	23.0	3.37	64.8	4.6	54.8	F
2	200.99	22.0	7.88	76.8	8.5	90.3	.
3	200.99	20.0	0.67	48.4	5.8	29.7	F
4	200.99	20.0	0.89	71.6	9.8	43.7	F

5 rows × 25 columns

Feature Scaling (StandardScaler)

In [8]:

```
scaler = StandardScaler()
df_scaled = df_encoded.copy()

df_scaled[numERIC_cols] = scaler.fit_transform(df_scaled[numERIC_cols])

df_scaled.head()
```

Out[8]:

	student_id	age	study_hours	class_attendance	sleep_hours	exam_score	gen
0	200.99	-1.520442	-0.531832	1.324086	0.225999	-0.191095	
1	200.99	1.106067	-0.276227	-0.301899	-1.390586	-0.407934	
2	200.99	0.668315	1.677629	0.392472	0.861085	1.469576	
3	200.99	-0.207188	-1.445942	-1.250872	-0.697764	-1.735413	
4	200.99	-0.207188	-1.350632	0.091578	1.611643	-0.994987	


5 rows × 25 columns



Histogram of a feature

In [9]:


```
plt.figure(figsize=(8,4))
plt.hist(df['study_hours'], bins=15)
plt.title("Distribution of Study Hours")
plt.xlabel("Study Hours")
plt.ylabel("Frequency")
plt.show()
```




Scatter plot to see relationships

In [10]:


```
plt.figure(figsize=(8,5))
sns.scatterplot(x="study_hours", y="exam_score", hue="gender", data=df)
plt.title("Study Hours vs Exam Score")
plt.show()
```



```
In [11]: plt.figure(figsize=(8,5))
sns.scatterplot(x="class_attendance", y="exam_score", hue="gender", data=df)
plt.title("Class Attendance vs Exam Score")
plt.show()
```




```
In [12]: plt.figure(figsize=(8,5))
sns.scatterplot(x="sleep_hours", y="study_hours", hue="gender", data=df)
plt.title("Sleep Hours vs Study Hours")
plt.show()
```




Correlation Heatmap

```
In [13]: plt.figure(figsize=(12,8))
sns.heatmap(df_encoded.corr(), annot=False, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```



```
In [14]: plt.figure(figsize=(8,5))
sns.heatmap(df[numERIC_COLS].corr(), annot=True, cmap="coolwarm")
plt.title("Correlation Heatmap (Numeric Only)")
plt.show()
```



In []: