

Multiple Linear Regression Experiment

Exam Score Prediction

This notebook implements a multiple linear regression model to predict exam scores based on various student factors.

```
In [18]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import warnings
warnings.filterwarnings('ignore')
```

1. Load the Dataset

```
In [19]: # Load the dataset
df = pd.read_csv('Exam_Score_Prediction.csv')
print("Dataset Shape:", df.shape)
print("\nFirst few rows:")
df.head()
```

Dataset Shape: (20000, 13)

First few rows:

```
Out[19]:
```

	student_id	age	gender	course	study_hours	class_attendance	internet_access	sle
0	200.99	17	male	diploma	2.78	92.9	yes	
1	200.99	23	other	bca	3.37	64.8	yes	
2	200.99	22	male	b.sc	7.88	76.8	yes	
3	200.99	20	other	diploma	0.67	48.4	yes	
4	200.99	20	female	diploma	0.89	71.6	yes	

2. Exploratory Data Analysis

```
In [20]: # Dataset information
print("Dataset Info:")
print(df.info())
print("\n" + "*50")
print("\nDescriptive Statistics:")
print(df.describe())
```

```
print("\n" + "="*50)
print("\nMissing Values:")
print(df.isnull().sum())
```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 13 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   student_id        20000 non-null   float64
 1   age                20000 non-null   int64  
 2   gender              20000 non-null   object  
 3   course              20000 non-null   object  
 4   study_hours         20000 non-null   float64
 5   class_attendance    20000 non-null   float64
 6   internet_access     20000 non-null   object  
 7   sleep_hours          20000 non-null   float64
 8   sleep_quality        20000 non-null   object  
 9   study_method         20000 non-null   object  
 10  facility_rating      20000 non-null   object  
 11  exam_difficulty      20000 non-null   object  
 12  exam_score           20000 non-null   float64
dtypes: float64(5), int64(1), object(7)
memory usage: 2.0+ MB
None
```

=====

Descriptive Statistics:

	student_id	age	study_hours	class_attendance	\
count	20000.000000	20000.000000	20000.000000	20000.000000	
mean	10000.500000	20.473300	4.007604	70.017365	
std	5770.211372	2.284458	2.308313	17.282262	
min	200.990000	17.000000	0.080000	40.600000	
25%	5000.750000	18.000000	2.000000	55.100000	
50%	10000.500000	20.000000	4.040000	69.900000	
75%	15000.250000	22.000000	6.000000	85.000000	
max	19800.010000	24.000000	7.910000	99.400000	

	sleep_hours	exam_score
count	20000.000000	20000.000000
mean	7.00856	62.513225
std	1.73209	18.908491
min	4.10000	19.599000
25%	5.50000	48.800000
50%	7.00000	62.600000
75%	8.50000	76.300000
max	9.90000	100.000000

=====

Missing Values:

student_id	0
age	0
gender	0
course	0
study_hours	0
class_attendance	0
internet_access	0
sleep_hours	0
sleep_quality	0
study_method	0
facility_rating	0

```
exam_difficulty      0
exam_score          0
dtype: int64
```


```
In [21]: # Check target variable distribution
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.hist(df['exam_score'], bins=30, edgecolor='black', color='skyblue')
plt.xlabel('Exam Score')
plt.ylabel('Frequency')
plt.title('Distribution of Exam Scores')
plt.grid(True, alpha=0.3)

plt.subplot(1, 2, 2)
plt.boxplot(df['exam_score'])
plt.ylabel('Exam Score')
plt.title('Box Plot of Exam Scores')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print(f"Mean Exam Score: {df['exam_score'].mean():.2f}")
print(f"Median Exam Score: {df['exam_score'].median():.2f}")
print(f"Std Deviation: {df['exam_score'].std():.2f}")
```



```
Mean Exam Score: 62.51
Median Exam Score: 62.60
Std Deviation: 18.91
```

3. Data Preprocessing

```
In [22]: # Identify categorical and numerical columns
categorical_cols = df.select_dtypes(include=['object']).columns.tolist()
numerical_cols = df.select_dtypes(include=['int64', 'float64']).columns.tolist()

print("Categorical Columns:", categorical_cols)
print("\nNumerical Columns:", numerical_cols)
```

```
Categorical Columns: ['gender', 'course', 'internet_access', 'sleep_quality', 'study_method', 'facility_rating', 'exam_difficulty']
```

```
Numerical Columns: ['student_id', 'age', 'study_hours', 'class_attendance', 'sleep_hours', 'exam_score']
```

```
In [23]: # Encode categorical variables
df_encoded = df.copy()
label_encoders = {}

for col in categorical_cols:
    le = LabelEncoder()
    df_encoded[col] = le.fit_transform(df[col])
    label_encoders[col] = le
    print(f"{col}: {dict(zip(le.classes_, le.transform(le.classes_)))}")

print("\nEncoded Dataset:")
df_encoded.head()

gender: {'female': np.int64(0), 'male': np.int64(1), 'other': np.int64(2)}
course: {'b.com': np.int64(0), 'b.sc': np.int64(1), 'b.tech': np.int64(2), 'ba': np.int64(3), 'bba': np.int64(4), 'bca': np.int64(5), 'diploma': np.int64(6)}
internet_access: {'no': np.int64(0), 'yes': np.int64(1)}
sleep_quality: {'average': np.int64(0), 'good': np.int64(1), 'poor': np.int64(2)}
study_method: {'coaching': np.int64(0), 'group study': np.int64(1), 'mixed': np.int64(2), 'online videos': np.int64(3), 'self-study': np.int64(4)}
facility_rating: {'high': np.int64(0), 'low': np.int64(1), 'medium': np.int64(2)}
exam_difficulty: {'easy': np.int64(0), 'hard': np.int64(1), 'moderate': np.int64(2)}


Encoded Dataset:
```

	student_id	age	gender	course	study_hours	class_attendance	internet_access	sleep_quality
0	200.99	17	1	6	2.78	92.9	1	1
1	200.99	23	2	5	3.37	64.8	1	1
2	200.99	22	1	1	7.88	76.8	1	1
3	200.99	20	2	6	0.67	48.4	1	1
4	200.99	20	0	6	0.89	71.6	1	1

4. Feature Correlation Analysis

```
In [24]: # Correlation matrix
plt.figure(figsize=(14, 10))
correlation_matrix = df_encoded.corr()
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm', center=0)
plt.title('Correlation Matrix of Features')
plt.tight_layout()
plt.show()

# Top correlations with exam_score
print("\nCorrelation with Exam Score:")
print(correlation_matrix['exam_score'].sort_values(ascending=False))
```



Correlation with Exam Score:

```

exam_score          1.000000
study_hours         0.717788
class_attendance   0.308850
sleep_hours         0.133222
age                 0.006522
exam_difficulty    0.003432
student_id          0.003267
gender              -0.000428
course              -0.000448
internet_access    -0.007826
facility_rating     -0.071879
sleep_quality       -0.099725
study_method        -0.154610
Name: exam_score, dtype: float64

```

5. Feature Selection and Data Splitting

```
In [25]: # Prepare features and target
X = df_encoded.drop(['exam_score', 'student_id'], axis=1) # Drop target and student ID
y = df_encoded['exam_score']

print("Features shape:", X.shape)
print("Target shape:", y.shape)
print("\nFeature columns:")
print(X.columns.tolist())
```

```
Features shape: (20000, 11)
Target shape: (20000,)
```

```
Feature columns:
['age', 'gender', 'course', 'study_hours', 'class_attendance', 'internet_access',
'sleep_hours', 'sleep_quality', 'study_method', 'facility_rating', 'exam_difficulty']
```

```
In [26]: # Split the dataset into training and testing sets (80-20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
print("Training set size:", X_train.shape[0])
print("Testing set size:", X_test.shape[0])
print(f"\nTraining set: {len(X_train)} samples ({len(X_train)}/{len(X)*100:.1f}%)")
print(f"Testing set: {len(X_test)} samples ({len(X_test)}/{len(X)*100:.1f}%)")
```

```
Training set size: 16000
Testing set size: 4000
```

```
Training set: 16000 samples (80.0%)
Testing set: 4000 samples (20.0%)
```

6. Build Multiple Linear Regression Model

```
In [27]: # Create and train the model
model = LinearRegression()
model.fit(X_train, y_train)

print("Model trained successfully!")
print("\nModel Coefficients:")
coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': model.coef_
}).sort_values('Coefficient', key=abs, ascending=False)

print(coef_df.to_string(index=False))
print(f"\nIntercept: {model.intercept_:.4f}")
```

```
Model trained successfully!
```

```
Model Coefficients:
      Feature   Coefficient
study_hours      5.881195
sleep_quality    -2.345414
study_method     -2.040250
facility_rating   -1.812587
sleep_hours       1.435547
class_attendance  0.342715
exam_difficulty   0.149939
internet_access   0.147043
gender            -0.036329
course             0.029321
age                0.007428
```


```
Intercept: 12.6640
```

```
In [28]: # Visualize feature importance (coefficients)
plt.figure(figsize=(12, 6))
coef_df_sorted = coef_df.sort_values('Coefficient')
```

```

plt.barh(coef_df_sorted['Feature'], coef_df_sorted['Coefficient'], color='steelblue')
plt.xlabel('Coefficient Value')
plt.ylabel('Features')
plt.title('Feature Importance (Regression Coefficients)')
plt.axvline(x=0, color='red', linestyle='--', linewidth=1)
plt.tight_layout()
plt.show()

```



7. Model Predictions

```

In [29]: # Make predictions
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

print("Predictions made successfully!")
print(f"\nSample predictions (first 10):")
comparison_df = pd.DataFrame({
    'Actual': y_test.values[:10],
    'Predicted': y_test_pred[:10],
    'Difference': y_test.values[:10] - y_test_pred[:10]
})
print(comparison_df)

```

Predictions made successfully!

Sample predictions (first 10):

	Actual	Predicted	Difference
0	31.1	35.989417	-4.889417
1	81.6	71.391764	10.208236
2	68.0	52.903226	15.096774
3	100.0	93.428419	6.571581
4	84.8	80.749395	4.050605
5	52.8	32.126485	20.673515
6	53.1	55.510023	-2.410023
7	81.0	67.908527	13.091473
8	72.6	64.550403	8.049597
9	25.4	30.249939	-4.849939

```


In [30]: # Scatter plot: Actual vs Predicted values
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_test_pred, alpha=0.6, edgecolors='k', s=80)

```

```

plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel('Actual Exam Scores', fontsize=12)
plt.ylabel('Predicted Exam Scores', fontsize=12)
plt.title('Actual vs Predicted Exam Scores - Test Set', fontsize=14, fontweight='bold')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```



8. Model Evaluation

```

In [31]: # Calculate evaluation metrics
train_mse = mean_squared_error(y_train, y_train_pred)
test_mse = mean_squared_error(y_test, y_test_pred)
train_rmse = np.sqrt(train_mse)
test_rmse = np.sqrt(test_mse)
train_mae = mean_absolute_error(y_train, y_train_pred)
test_mae = mean_absolute_error(y_test, y_test_pred)
train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

print("*"*60)
print("MODEL EVALUATION METRICS")
print("*"*60)
print("\nTraining Set Performance:")
print(f" Mean Squared Error (MSE): {train_mse:.4f}")
print(f" Root Mean Squared Error (RMSE): {train_rmse:.4f}")
print(f" Mean Absolute Error (MAE): {train_mae:.4f}")
print(f" R2 Score: {train_r2:.4f}")

print("\nTesting Set Performance:")
print(f" Mean Squared Error (MSE): {test_mse:.4f}")
print(f" Root Mean Squared Error (RMSE): {test_rmse:.4f}")
print(f" Mean Absolute Error (MAE): {test_mae:.4f}")
print(f" R2 Score: {test_r2:.4f}")
print("*"*60)

```

```

# Check for overfitting/underfitting
print(f"\nR² Difference (Train - Test): {train_r2 - test_r2:.4f}")
if abs(train_r2 - test_r2) < 0.05:
    print("✓ Model is well-balanced (no significant overfitting)")
elif train_r2 > test_r2:
    print("⚠ Model may be slightly overfitting")
else:
    print("⚠ Model may be underfitting")

```

=====

MODEL EVALUATION METRICS

=====

Training Set Performance:

Mean Squared Error (MSE): 118.1183
Root Mean Squared Error (RMSE): 10.8682
Mean Absolute Error (MAE): 8.7191
R² Score: 0.6696

Testing Set Performance:

Mean Squared Error (MSE): 118.9531
Root Mean Squared Error (RMSE): 10.9066
Mean Absolute Error (MAE): 8.8006
R² Score: 0.6674

=====

R² Difference (Train - Test): 0.0021

✓ Model is well-balanced (no significant overfitting)

9. Visualization of Results

In [32]:

```

# Actual vs Predicted values
plt.figure(figsize=(15, 5))

# Training set
plt.subplot(1, 3, 1)
plt.scatter(y_train, y_train_pred, alpha=0.5, color='blue')
plt.plot([y_train.min(), y_train.max()], [y_train.min(), y_train.max()], 'r--',
plt.xlabel('Actual Scores')
plt.ylabel('Predicted Scores')
plt.title(f'Training Set\nR² = {train_r2:.4f}')
plt.grid(True, alpha=0.3)

# Testing set
plt.subplot(1, 3, 2)
plt.scatter(y_test, y_test_pred, alpha=0.5, color='green')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2
plt.xlabel('Actual Scores')
plt.ylabel('Predicted Scores')
plt.title(f'Testing Set\nR² = {test_r2:.4f}')
plt.grid(True, alpha=0.3)

# Combined view
plt.subplot(1, 3, 3)
plt.scatter(y_train, y_train_pred, alpha=0.3, color='blue', label='Training')
plt.scatter(y_test, y_test_pred, alpha=0.5, color='green', label='Testing')
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', lw=2, label='Perfect Pre
plt.xlabel('Actual Scores')
plt.ylabel('Predicted Scores')


```

```

plt.title('Combined View')
plt.legend()
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



```

In [33]: # Residual analysis
residuals_train = y_train - y_train_pred
residuals_test = y_test - y_test_pred

plt.figure(figsize=(15, 5))

# Residual plot for training
plt.subplot(1, 3, 1)
plt.scatter(y_train_pred, residuals_train, alpha=0.5, color='blue')
plt.axhline(y=0, color='r', linestyle='--', lw=2)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot - Training Set')
plt.grid(True, alpha=0.3)


# Residual plot for testing
plt.subplot(1, 3, 2)
plt.scatter(y_test_pred, residuals_test, alpha=0.5, color='green')
plt.axhline(y=0, color='r', linestyle='--', lw=2)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot - Testing Set')
plt.grid(True, alpha=0.3)

# Residual distribution
plt.subplot(1, 3, 3)
plt.hist(residuals_test, bins=30, edgecolor='black', color='lightgreen', alpha=0.8)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Distribution of Residuals (Test Set)')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print(f"Mean of residuals (should be close to 0): {residuals_test.mean():.4f}")
print(f"Std of residuals: {residuals_test.std():.4f}")

```




Mean of residuals (should be close to 0): -0.0394
 Std of residuals: 10.9079

```
In [34]: # Error distribution comparison
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
errors_test = np.abs(y_test - y_test_pred)
plt.hist(errors_test, bins=30, edgecolor='black', color='coral', alpha=0.7)
plt.xlabel('Absolute Error')
plt.ylabel('Frequency')
plt.title('Distribution of Absolute Errors (Test Set)')
plt.axvline(test_mae, color='red', linestyle='--', linewidth=2, label=f'MAE: {test_mae:.2f}')
plt.legend()
plt.grid(True, alpha=0.3)

plt.subplot(1, 2, 2)
plt.boxplot([residuals_train, residuals_test], labels=['Training', 'Testing'])
plt.ylabel('Residuals')
plt.title('Residuals Comparison')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



10. Model Summary

```
In [35]: # Model equation
print("=*80")
print("MULTIPLE LINEAR REGRESSION MODEL EQUATION")
print("=*80")
print(f"\nexam_score = {model.intercept_:.4f}", end="")
```

```

for feature, coef in zip(X.columns, model.coef_):
    sign = "+" if coef >= 0 else "-"
    print(f" {sign} {abs(coef):.4f} * {feature}", end="")
print("\n" + "="*80)

# Summary statistics
summary_data = {
    'Metric': ['R² Score', 'RMSE', 'MAE', 'MSE'],
    'Training': [f'{train_r2:.4f}', f'{train_rmse:.4f}', f'{train_mae:.4f}', f'{train_mse:.4f}'],
    'Testing': [f'{test_r2:.4f}', f'{test_rmse:.4f}', f'{test_mae:.4f}', f'{test_mse:.4f}']
}
summary_df = pd.DataFrame(summary_data)

print("\n\nMODEL PERFORMANCE SUMMARY:")
print(summary_df.to_string(index=False))
print("\n" + "="*80)

```

=====

MULTIPLE LINEAR REGRESSION MODEL EQUATION

=====

```

exam_score = 12.6640 + 0.0074 * age - 0.0363 * gender + 0.0293 * course + 5.8812
* study_hours + 0.3427 * class_attendance + 0.1470 * internet_access + 1.4355 * sleep_hours
- 2.3454 * sleep_quality - 2.0403 * study_method - 1.8126 * facility_rating
+ 0.1499 * exam_difficulty
=====
```

MODEL PERFORMANCE SUMMARY:

	Metric	Training	Testing
R ²	Score	0.6696	0.6674
RMSE		10.8682	10.9066
MAE		8.7191	8.8006
MSE		118.1183	118.9531

=====

11. Conclusions

Key Findings:

1. The multiple linear regression model has been successfully trained on the exam score prediction dataset
2. Model performance can be evaluated using R² score, RMSE, MAE, and MSE metrics
3. Feature importance is determined by the magnitude of regression coefficients
4. Residual analysis helps identify potential issues with the model assumptions